

A Tutorial on Machine Learning Methods

Varun Chandola

Computer Science and Engineering
Computational and Data-Enabled Science and Engineering
University at Buffalo, State University of New York

March 24, 2016

Tutorial Overview

Part I - Methods (Chandola)

- Today from 1.00 PM - 3.00 PM
- 330 Student Union
- Introduction and general overview of methods and concepts

Part II - Applications (Bauman and Hachmann)

- Tomorrow from 9.00 AM - 11.00 AM
- 280 Park Hall
- II.a - Bayesian Methods and Partial Differential Equations
- II.b - Machine Learning for Computational Chemistry

Setup

- Python and iPython Notebooks
- `scikit-learn` - ML library in Python

Follow Along

- Slides: <https://github.com/ubdsgroup/cdsedaystutorial/blob/master/talk.pdf>
- Notebook: <http://nbviewer.jupyter.org/github/ubdsgroup/cdsedaystutorial/blob/master/MachineLearningBasics.ipynb>
- Git Repo:
<https://github.com/ubdsgroup/cdsedaystutorial.git>

What makes a machine intelligent?

- 1 Talk. See. Hear.
 - Natural Language Processing, Computer Vision, Speech Recognition
- 2 Store. Access. Represent. (*Knowledge*)
 - Ontologies. Semantic Networks. Information Retrieval.
- 3 Reason.
 - Mathematical Logic. Bayesian Inference.
- 4 **Learn.**
 - Improve with Experience

What makes a machine intelligent?

- 1 Talk. See. Hear.
 - Natural Language Processing, Computer Vision, Speech Recognition
- 2 Store. Access. Represent. (*Knowledge*)
 - Ontologies. Semantic Networks. Information Retrieval.
- 3 Reason.
 - Mathematical Logic. Bayesian Inference.
- 4 **Learn.**
 - Improve with Experience
 - Machine Learning

What is Machine Learning?

- Computers learn without being **explicitly programmed**.
 - Arthur Samuel (1959)
- A computer program learns from experience E with respect to some task T , if its performance P while performing task T improves over E .
 - Tom Mitchell (1989)

Why Machine Learning?

- Machines that know everything from the beginning?
 - Too bulky. Creator already knows everything. Fails with *new* experiences.
- Machines that *learn*?
 - Compact. Learn what is *necessary*.
 - Adapt.
 - Assumption: Future experiences are not too different from past experiences.
 - Have (structural) relationship.



Why is ML so popular?


- Learn to do tasks that are “impossible” (or infeasible) for humans
- Automatically learn from past to operate in the unknown future
- In CS, we focus on issues such as scalability, usability, performance, error handling, complexity, convergence, *big data challenges*
- What is in it for non-AI and non-CS folks?
 - All of us are writing tiny “robots”
 - Almost all machine learning \equiv learning from **data**

Topics in Machine Learning

- Tasks
 - **Classification, Regression**, Ranking, Latent Variable Modeling (**clustering, dimensionality reduction**, dictionary learning)
- Learning types
 - **Supervised, Unsupervised**, Semi-supervised, Reinforcement
- Learning strategies
 - **Error minimization**, Statistical, **Heuristic search**
- Inference models
 - Parametric: **Probabilistic, Linear, Non-linear**
 - Non-parametric
- Cross-cutting issues
 - Handling structured inputs/outputs, **Stability/Generalizability, Interpretability**, Incorporating domain knowledge, Scalability

What else will I not cover?

- What is Deep Learning?
- How does AlphaGo work?
 - It does use ML (neural networks to be exact)
 - <https://www.tastehit.com/blog/google-deepmind-alphago-how-it-works/>
- How to make money in the stock market with this?

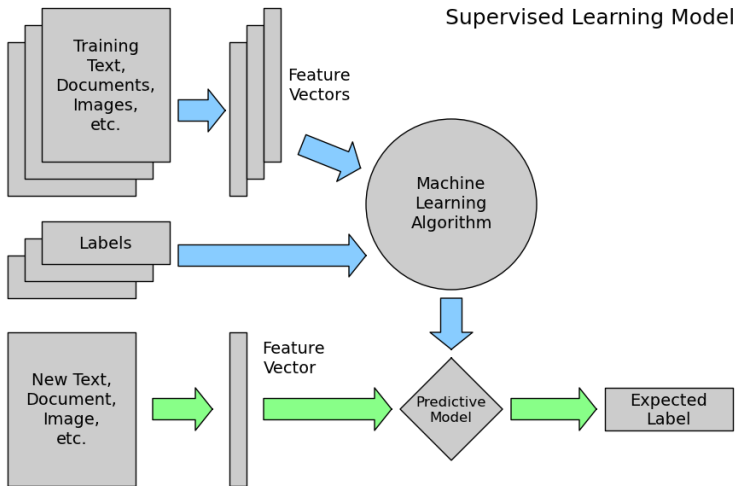


Src: Nature, 2016

Classification or Regression



- Given input \mathbf{x} , predict target \mathbf{y}
- Assumptions:
 - 1 $\mathbf{x} \in \mathbb{R}^D$, i.e., \mathbf{x} is represented as a vector of D real values – attributes or features
 - 2 Regression: $y \in \mathbb{R}$
 - 3 Classification $y \in \{+1, -1\}$ or $y \in \{1, 2, \dots, k\}$

A Data Science View





Parametric vs. Non-parametric Models

Parametric

- Model represented as a set of *parameters*
- Concise, easy to store 
- Making assumptions (inductive bias) 
- *Examples:* Linear regression, Neural networks

Non-parametric

- No (finite) set of parameters
- Use entire training data for inference on a new example
- No assumptions 
- Complexity grows with data 
- *Examples:* Gaussian process regression, Nearest neighbor classification

Supervised Learning for Parametric Models

- Given a training data set, learn (or estimate) the model parameters

Inverse Problem Theory - Physical Systems

- Forward modeling: Given system parameters one can compute the behavioral output of the system (apply physical laws)
- Inverse modeling: Given the observations, *infer* the values of the system parameters

Possible approaches

- Find parameters that minimize error on training data
 - Least squares regression, neural networks, support vector machines
- Find parameters that maximize the likelihood (or posterior) of the data distribution
 - Naive Bayes, Bayesian Regression
- Search for best solution using greedy heuristics
 - Decision trees, Random Forests

Error Minimization Methods



- 1 Define an error as a function of the model parameters (Θ) and the training data
- 2 Minimize the error to obtain optimal parameters

$$\begin{aligned}\hat{\Theta} &= \arg \min_{\Theta} J(\Theta) \\ &= \arg \min_{\Theta} \sum_{i=1}^N (y_i - \hat{y}_i)^2\end{aligned}$$

- The squared error between true output (y_i) and predicted output (\hat{y}_i) is not the only choice

Perils of Unconstrained Error Minimization Approaches

- Favor complex solutions
- Do perfectly on training data (overfitting)
- No guarantee to work well on unseen data (generalizability)
- Learning theory: Complex models \Rightarrow Poor generalizability



How to Control Overfitting?

- Use simpler models (linear instead of polynomial)
 - Might have poor results (underfitting)
- Use regularized complex models

$$\hat{\Theta} = \arg \min_{\Theta} J(\Theta) + \alpha R(\Theta)$$

- $R()$ corresponds to the penalty paid for complexity of the model

More about Regularization

- **Probabilistic methods:** impose regularization through priors on the parameters
- **Search methods:** avoid searching along paths that lead to complex results
- Some regularization methods have other benefits too:
 - Feature selection
 - Imposing sparsity constraints (interpretability)
 - Incorporating domain knowledge
 - Reducing impact of correlated inputs.

Regularization for Linear Models

- Linear regression (for regression) and logistic regression (for classification) are two popular linear models

$$y \sim \mathcal{N}(\mathbf{w}^\top \mathbf{x}, \sigma^2)$$

$$y \sim \text{Bernoulli}(\text{sigmoid}(\mathbf{w}^\top \mathbf{x}))$$

- where $\text{sigmoid}(a) = \frac{1}{1 + \exp(-a)}$
- One can use polynomial expansion to model non-linear dependencies (the model is still linear in the weights)
- How to control the complexity?

Examples of Regularization

Ridge Regression

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} J(\mathbf{w}) + \alpha \|\mathbf{w}\|^2$$

- Also known as l_2 or *Tikhonov* regularization
- Helps in reducing impact of correlated inputs

LASSO

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} J(\mathbf{w}) + \alpha |\mathbf{w}|$$

- Also known as l_1 regularization
- Helps in feature selection – favors sparse solutions

Regularization – A Research Area by Itself

- For linear models, many other forms exist
 - 1 Elastic net
 - 2 Group Lasso
 - 3 Tree Structured Lasso
 - 4 Network Lasso
- Each formulation requires a unique optimization problem to be solved


Line as a Decision Surface

- Decision boundary represented by the hyperplane \mathbf{w}
- For binary classification, \mathbf{w} points **towards** the positive class

Decision Rule

$$y = \text{sign}(\mathbf{w}^\top \mathbf{x} + b)$$

- $\mathbf{w}^\top \mathbf{x} + b > 0 \Rightarrow y = +1$
- $\mathbf{w}^\top \mathbf{x} + b < 0 \Rightarrow y = -1$



Support Vector Machines

- A hyperplane based classifier defined by \mathbf{w} and b
- Find hyperplane with *maximum separation margin* on the training data

SVM Prediction Rule

$$y = \text{sign}(\mathbf{w}^\top \mathbf{x} + b)$$

SVM Learning

- **Input:** Training data $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$
- **Objective:** Learn \mathbf{w} and b that maximizes the margin

Going Non-Linear with SVMs

- SVMs are linear classifiers
- To model non-linear dependencies
 - 1 Basis vector expansion

$$\Phi(x) = 1, x, x^2, \dots, x^p$$

still restrictive

2 The **kernel trick**


- Allows for (implicit) mapping of data into a new high dimensional space through kernel functions

$$k(x_1, x_2) \equiv \Phi_1^\top \Phi_2$$

- The classes are assumed to be linearly separable in that space
- Most popular is Radial Basis Function (RBF) kernel that allows mapping data into infinite dimensional space!

Decision Trees

- Learning algorithm searches for a sequence of rules that organize a data set
- Learning is one variable at a time
 - Find boundaries between sections of data (a sub-decision)
- Which variable to choose and where to set the boundary?
 - Information Gain, Gini Impurity, Variance Reduction



Why Doesn't Everyone Use Decision Trees?

Strengths

- Both regression and classification
- Handle all types of attributes, missing data, noise
- Validate through statistical tests
- Interpretable
- No need for normalization

Weaknesses

- No optimality guarantee (learning problem is NP-complete)
- Favor attributes with more levels or categories
- Become too complex for certain types of problems


Extensions of Decision Trees

- Using conditional inference
- Regression trees
- Boosted Trees
- Random Forests


What is Clustering?

- Grouping similar things together
- A notion of a similarity or distance metric
- A type of **unsupervised learning**
 - Learning without any labels or target

Expected Outcome of Clustering



Expected Outcome of Clustering



K-Means Clustering

- **Objective:** Group a set of N points ($\in \mathbb{R}^D$) into K clusters.

K-Means Clustering

- **Objective:** Group a set of N points ($\in \mathbb{R}^D$) into K clusters.
- 1 **Start** with k *randomly initialized* points in D dimensional space
 - Denoted as $\{\mathbf{c}_k\}_{k=1}^K$
 - Also called *cluster centers*

K-Means Clustering

- **Objective:** Group a set of N points ($\in \mathbb{R}^D$) into K clusters.
- ① **Start** with k *randomly initialized* points in D dimensional space
 - Denoted as $\{\mathbf{c}_k\}_{k=1}^K$
 - Also called *cluster centers*
- ② **Assign** each input point \mathbf{x}_n ($\forall n \in [1, N]$) to cluster k , such that:

$$\min_k \text{dist}(\mathbf{x}_n, \mathbf{c}_k)$$

K-Means Clustering

- **Objective:** Group a set of N points ($\in \mathbb{R}^D$) into K clusters.
- ① **Start** with k *randomly initialized* points in D dimensional space
 - Denoted as $\{\mathbf{c}_k\}_{k=1}^K$
 - Also called *cluster centers*

- ② **Assign** each input point \mathbf{x}_n ($\forall n \in [1, N]$) to cluster k , such that:

$$\min_k \text{dist}(\mathbf{x}_n, \mathbf{c}_k)$$

- ③ **Revise** each cluster center \mathbf{c}_k using all points assigned to cluster k

K-Means Clustering

- **Objective:** Group a set of N points ($\in \mathbb{R}^D$) into K clusters.
- 1 **Start** with k *randomly initialized* points in D dimensional space
 - Denoted as $\{\mathbf{c}_k\}_{k=1}^K$
 - Also called *cluster centers*

- 2 **Assign** each input point \mathbf{x}_n ($\forall n \in [1, N]$) to cluster k , such that:

$$\min_k \text{dist}(\mathbf{x}_n, \mathbf{c}_k)$$

- 3 **Revise** each cluster center \mathbf{c}_k using all points assigned to cluster k
- 4 **Repeat** 2

Variants of K-Means

- Finding distance
 - Euclidean distance is popular
- Finding cluster centers
 - Mean for K-Means
 - Median for k-medoids

Choosing Parameters

- 1 Similarity/distance metric
 - Can use non-linear transformations
 - K-Means with Euclidean distance produces “circular” clusters
- 2 How to set k ?
 - Trial and error
 - How to evaluate clustering?
 - K-Means objective function

$$J(\mathbf{c}, \mathbf{R}) = \sum_{n=1}^N \sum_{k=1}^K R_{nk} \|\mathbf{x}_n - \mathbf{c}_k\|^2$$

- \mathbf{R} is the cluster assignment matrix

$$R_{nk} = \begin{cases} 1 & \text{If } \mathbf{x}_n \in \text{cluster } k \\ 0 & \text{Otherwise} \end{cases}$$

Initialization Issues

- Can lead to wrong clustering
- Better strategies
 - 1 Choose first centroid randomly, choose second farthest away from first, third farthest away from first and second, and so on.
 - 2 Make multiple runs and choose the best

Strengths and Limitations of K-Means

Strengths

- Simple
- Can be extended to other types of data
- Easy to parallelize

Weaknesses

- Circular clusters (not with kernelized versions)
- Choosing K is always an issue
- Not guaranteed to be optimal
- Works well if natural clusters are round and of equal densities
- **Hard Clustering**

Issues with K-Means

- “Hard clustering”
- Assign every data point to exactly one cluster
- **Probabilistic Clustering**
 - Each data point can belong to multiple clusters with varying probabilities
 - In general

$$P(\mathbf{x}_i \in C_j) > 0 \quad \forall j = 1 \dots K$$

- For hard clustering probability will be 1 for one cluster and 0 for all others

Latent Variable Models

- Consider a probability distribution parameterized by θ
- Generates samples (\mathbf{x}) with probability $p(\mathbf{x}|\theta)$

2-step generative process

Latent Variable Models

- Consider a probability distribution parameterized by θ
- Generates samples (\mathbf{x}) with probability $p(\mathbf{x}|\theta)$

2-step generative process

- 1 Distribution generates the hidden variable

Latent Variable Models

- Consider a probability distribution parameterized by θ
- Generates samples (\mathbf{x}) with probability $p(\mathbf{x}|\theta)$

2-step generative process

- 1 Distribution generates the hidden variable
- 2 Distribution generates the observation, given the hidden variable

More About Latent Variable Models

- The observed random variable \mathbf{x} depends on a hidden random variable \mathbf{z}
- \mathbf{z} is generated using a *prior* distribution - $p(\mathbf{z})$
- \mathbf{x} is generated using $p(\mathbf{x}|\mathbf{z})$
- Different combinations of $p(\mathbf{z})$ and $p(\mathbf{x}|\mathbf{z})$ give different latent variable models
 - 1 Mixture Models
 - 2 Factor analysis
 - 3 Probabilistic Principal Component Analysis (PCA)
 - 4 Latent Dirichlet Allocation (LDA)
 - 5 Dictionary Learning, Sparse Coding

Other Resources

- My UB course - CSE474/574 (<http://www.cse.buffalo.edu/~chandola/machinelearning.html>)
- Beautiful ML (<http://www.r2d3.us/visual-intro-to-machine-learning-part-1/>)